




MODELISE

THE GOVERNANCE LAYER FOR THE AI

PROBLEM: AI IS EVERYWHERE. GOVERNANCE IS NOWHERE

- **The "Blank Check" Problem:** Unmonitored API usage leads to massive token waste and unpredictable costs.
- **Compliance Paralysis:** Government and mid tier firms can't use powerful models because they can't control *how* employees interact with them.
- **Wasted Resources:** Developers are wasting months building custom security layers from scratch instead of actually scaling their AI.
- **AI Misuse:** Stop AI from being weaponized or misused by wrapping every model in a strict, customizable policy engine.

TARGET SEGMENTS:

Four segments, one platform.

01 AI Platforms

HuggingFace, Replicate, Together AI

Offer governance as a hosted add on. Help platform builders give their users safety controls without building infra from scratch.

02 Enterprises

Mid-tier companies

Control what AI models employees can do. Avoid token waste, data leaks, and policy violations — without a dedicated AI safety team.

03 Government

Public sector, regulated agencies

Unlock AI adoption where it has been blocked by compliance concerns. Auditable, configurable, explainable governance out of the box.

04 Indie Developers

Solo builders, small teams

Ship AI powered products with governance baked in — no infra burden. Compete with enterprise grade safety on a startup budget.

The Regulatory Gold Rush


- *"As AI laws get stricter, the 'Governance Layer' becomes the most valuable piece of the AI stack."*
- **Regulatory-as-a-Service:** We allow companies to instantly update their AI behavior to match new laws in the EU, US, or India with a single toggle.
- **Market Opportunity:** The AI Governance market is projected to grow **48% annually**, reaching nearly **\$7B by 2030**.
- **Sovereign AI:** Government offices and highly regulated sectors can finally use LLMs because MODELISE provides the "Air Gap" of safety they require.

APPROACH

- **Plug-and-Play Governance Proxy:** Sits seamlessly between user prompts and AI models, automatically wrapping any model with enterprise-grade security without requiring developers to build custom guardrails.
- **Universal Deployment Flexibility:** Works out of the box for custom self-hosted ML models (Text & Vision) or external commercial APIs like Claude, Gemini, and OpenAI.
- **Layer 1 — Instant Multimodal Security:** Automatically filters prompt injections, malicious inputs, broken code, and unsafe text/image content before reaching your core model.
- **Layer 2 — Custom Policy Engine (Our Core MOAT):** Empowers organizations to upload custom compliance rules, corporate policies, and sensitive data redaction guidelines.
- **Geo-Specific & Regional Compliance Rules:** Allows country-specific regulations to automatically censor locally sensitive information and ful-fill regional AI laws (e.g., EU AI Act, US, or India).
- **Seamless Layer Chaining:** Integrates Layer 1 security with Layer 2 custom/country policies into a single automated pipeline, generating full audit trails for every prompt. Could use them separately too.
- **Zero-Code Prompt Inspection:** Functions as a unified proxy that intercepts and inspects prompts prior to model execution, preventing token waste and data leaks.

SOLUTION:

A plug and play governance layer for any AI model.



- Customizable policies per org, team, or use case
- Risk based enforcement : Redact sensitive info, escalate suspicious prompts.
- Real time monitoring of inputs & outputs
- Works with any model : no lock in
- Intelligent Policy Enforcement: Move beyond "Yes/No" blocking.

THE VISION

- **Vision:** Provide a zero trust governance layer that eliminates financial token-burn attacks and compliance violations before prompts reach execution models.
- **100% Financial Exploit Immunity:** Neutralizes token-burn and recursive prompt-drain attacks before downstream API costs are incurred.
- **Real-Time Policy Enforcement:** Automates compliance and data redaction dynamically without needing custom code per app or dedicated safety teams.
- **Air Gapped Sovereign AI:** Keeps evaluations fully on-premise, enabling safe AI adoption for government and regulated enterprises.




Technical Stack:

- **Framework & API Gateway:** FastAPI (Python 3.11+) with Pydantic v2 for async REST endpoints.
- **Vector Search & Policy RAG:** FAISS (faiss-cpu) using sentence-transformers/all-MiniLM-L6-v2 for normalized L_2 distance queries.
- **Local Inspection Engine:** Ollama / vLLM serving gemma2:2b for intent and cost analysis.
- **Security & Auth:** OAuth2 with JWT tokens (PyJWT) and Bcrypt password hashing (passlib).
- **Document Processing:** LangChain utilities (PyPDFLoader, RecursiveCharacterTextSplitter) for async chunking and ingestion.
- **Concurrency:** asyncio for non-blocking local model evaluation and multi-chunk FAISS searches.
- Training own DL models. Docker.

Secure & Trustworthy Architecture

- **One-Time On-Prem Hardware Setup:** Runs on cheap, local commodity hardware or edge devices (e.g., standard PC, consumer GPU, or Raspberry Pi) using lightweight models like gemma2:2b, eliminating recurring third-party inspection API fees.
- **Integrated Layer 1 Defense:** Built directly downstream of Layer 1 baseline security filters, forming a unified two-tiered firewall that neutralizes prompt injections before policy evaluation begins.
- **Complete Air-Gapped Data Privacy:** All evaluations happen locally within your private network—zero prompt data is transmitted to external vendors.
- **Dual-Gate Short-Circuiting:** Gate 1 (Cost Guard) intercepts recursive loops and budget overspends upfront, while Gate 2 (Policy RAG) halts execution at the first policy chunk violation without triggering the primary LLM.
- **Admin-Authenticated Policy Management:** All policy uploads, vector index updates, and threshold modifications are strictly locked behind JWT and Bcrypt admin password authentication.
- **Zero-Trust Budgeting & Anomaly Cutoffs:** Tracks live token spend per user_id with sliding-window rate limiters to automatically kill malicious traffic spikes.
- **Dynamic In-Memory Rule Updates:** Admins can hot-reload PDF/CSV compliance rules directly into active vector memory without server restarts



THANK YOU

- Team Name: CyberHawk
- Team Code: SCRY-D8C843
- ATUL KUMAR MAURYA
- ANSHAM MAURYA